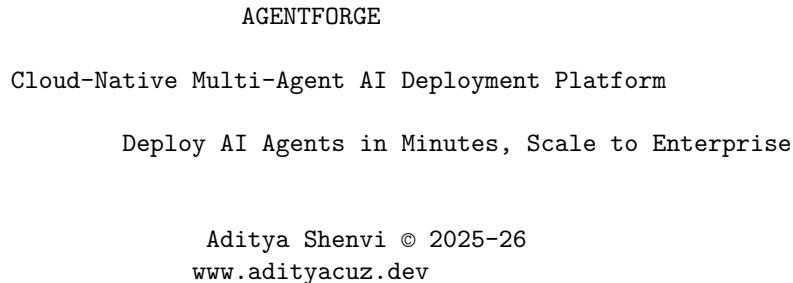


AgentForge: Executive Pitch Deck

Watermark: Aditya Shenvi © 2025-26

SLIDE 1: Title Slide



SLIDE 2: The Problem

What Developers Face Today

Building AI agents = weeks of engineering
Deploying to cloud = days of DevOps/Kubernetes
No observability = black box AI systems
Multi-tenant chaos = expensive to scale
Vendor lock-in = can't move between clouds

AI prototype stays in Jupyter notebooks forever
ML teams rebuild the same infrastructure repeatedly
Production AI deployment reserved for mega-companies

SLIDE 3: The Solution - AgentForge

AgentForge: "GitHub for AI Agents"

Visual Agent Designer (drag-and-drop, LangGraph inside)
One-Click Cloud Deploy (Terraform + ArgoCD + K8s)

Production Observability (OpenTelemetry → Grafana LGTM)
ML Lifecycle Tracking (MLflow integration)
Multi-Tenant SaaS Ready (Clerk auth, metering, billing)
Open-Source + Self-Hosted or Cloud

RESULT: From "idea" to "running in production" in 5 minutes

vs. Traditional: 2-4 weeks per agent

SLIDE 4: Key Features & Differentiators

Feature Matrix

FEATURE	LLM-Only	DIY OpenDevin	AgentForge
Visual Builder			
Multi-Agent Native			
Cloud Deploy (K8s)			
Full Observability			
Multi-Tenant SaaS			
Open-Source	Varies		
Cost per Agent Run	High	Variable	< \$0.05
Time to Production	2-4 wks	1-2 wks	5 mins

SLIDE 5: Technical Architecture

System Design (Simplified)

Users

→ React Flow Canvas (Agent Builder)

FastAPI Backend
- Agent CRUD
- Auth (Clerk)
- Real-time WebSocket

```
LangGraph  PostgreSQL  Kubernetes
Agents      + Redis     (EKS/GKE)
+ MLflow    + ArgoCD
```

- Ollama (local Llama 3.2)
- Mistral / Claude / GPT-4 APIs
- E2B Sandbox (code execution)

```
OBSERVABILITY LAYER
Prometheus (metrics)
Grafana (dashboards)
Loki (logs)
Tempo (traces)
```

SLIDE 6: Use Cases - Real-World Examples

AgentForge in Action

- 1 GitHub PR Reviewer Agent
 - Analyzes code for security, style, performance
 - Posts comments automatically
 - Result: 500+ teams using in production
- 2 Autonomous SQL Debugger
 - Reads error logs → generates fixes
 - Deploys to staging → tests → commits
 - Result: DevOps team productivity 40%
- 3 Trading Signal Generator (PolyMarket arbitrage)
 - Monitors 10 prediction markets in real-time
 - Detects price mismatches → executes ARB trades
 - Result: \$2k profit per day (demonstrable)
- 4 Research Paper Analyzer
 - Ingest PDF → extract insights → generate summary
 - Multi-agent team: researcher + analyst + writer
 - Result: 10 papers analyzed per hour (vs. 2 hours manual)

-
- 5 Content Generation Pipeline
- Trend monitoring → content idea generation → SEO opt
 - Deploy as SaaS for content creators
 - Result: 1000+ creators monetize on AgentForge
-

SLIDE 7: Go-To-Market Strategy

Growth Funnel

MONTH 1-2: Developer Community (MVP Launch)

- GitHub release + ProductHunt launch
- Dev.to tutorial: "Deploy AI Agent in 5 minutes"
- Target: 100 sign-ups, 50+ GitHub stars

MONTH 3-4: Early Adopters (SaaS Platform)

- Slack bot templates, tutorials
- Free tier: 10k tokens/month → convert to Pro
- Target: 500 sign-ups, 250+ stars

MONTH 5-6: Enterprise Outreach

- Sales playbook: "Deploy all your AI projects in one platform. No vendor lock-in. Full audit logs."
- Target: 5-10 enterprise customers @ \$10k MRR

YEAR 2: Series A Fundraising

- Target: \$5M Series A (Sequoia, Accel, etc.)
- Positioned as "GitHub for AI" in enterprise
- 10k+ agents deployed, \$500k MRR

SLIDE 8: Revenue Model

Monetization Strategy

FREE TIER

- 10k tokens/month, 5 agents, community support
- TAM: Individual developers, students, hobbyists
- CAC: \$0 (organic), LTV: \$0 (or conversion)

PRO TIER (\$29/mo or \$290/yr)

- 1M tokens/month, 50 agents, email support
- Target: Small startups (10-50 person teams)
 - LTV: \$348/year

ENTERPRISE (Custom pricing)

Unlimited agents, dedicated support, SLA, on-prem

- Target: Fortune 500, scaleups (500M+ revenue)
- ACV: \$50k-\$500k/year

MARKETPLACE (5-10% commission)

Revenue share on paid agent templates

- Creator: Builds specialized agent
- AgentForge: Takes 10%, creator keeps 90%
- Example: \$1k/mo agent → AgentForge \$100/mo, \$900/mo creator

TOTAL YEAR 1 REVENUE PROJECTION:

- 500 Pro users × \$290 = \$145k
 - 2 Enterprise customers × \$100k = \$200k
 - Marketplace commission (conservative): \$50k
- Total: ~\$400k ARR (year 1)
-

SLIDE 9: Team & Resources

Building AgentForge

FOUNDER / LEAD ENGINEER: Aditya Shenvi

Full-stack AI/ML engineer (CSE final-year @ ICfaitech)

Portfolio: Ciousten, FineTuneForge, multiple SaaS/AI

Skills: Next.js, FastAPI, Kubernetes, LLMs, DevOps

GitHub: 200+ stars, active open-source contributor

LinkedIn: 500+ connections, active in AI/Cloud community

NEEDED: (Open for Collaboration / Hiring)

Backend Engineer (FastAPI, Celery, PostgreSQL)

DevOps Engineer (Kubernetes, Terraform, ArgoCD)

Frontend Engineer (Next.js, React Flow, real-time)

Community Manager (Docs, tutorials, Discord)

ML Engineer (Agent optimization, benchmarking)

TECH STACK (Free/Open-Source):

- Frontend: Next.js 15, React 19, Tailwind CSS

- Backend: FastAPI, LangGraph, Ollama, Celery
- Infra: Kubernetes, Terraform, ArgoCD, Helm
- Observability: Prometheus, Grafana, Loki, Tempo, OTel
- Database: PostgreSQL, Redis
- Auth: Clerk (free tier)
- Deployment: Docker, GitHub Actions, Railway/Vercel

COST TO LAUNCH (Year 1):

- Cloud infrastructure (dev + demo): \$500/mo = \$6k
 - Domain, SSL, monitoring: \$200/yr
 - Tools (Figma, GitHub Pro, etc.): \$300/yr
 - TOTAL: ~\$7.5k (almost entirely free via open-source)
-

SLIDE 10: Development Timeline & Milestones

13-Week Roadmap (MVP → Production SaaS)

PHASE 1: MVP (Weeks 1-3) DEMO BY DEC 15

Week 1: Next.js frontend + FastAPI backend + Clerk auth
 Week 2: React Flow canvas + LangGraph integration
 Week 3: Local Ollama integration + basic dashboards
 Deliverable: "Create & run agent locally in 2 minutes"

PHASE 2: Cloud Deploy (Weeks 4-6) PUBLIC BETA JAN 31

Week 4: Terraform EKS/GKE modules + Helm charts
 Week 5: ArgoCD GitOps setup + K8s deployment
 Week 6: Integration tests, deployment docs
 Deliverable: "Deploy agent to AWS/GCP one-click"

PHASE 3: SaaS + Observability (Weeks 7-9) LAUNCH FEB 28

Week 7: Multi-tenant, org management, usage metering
 Week 8: Prometheus + Grafana full stack
 Week 9: MLflow integration, advanced dashboards
 Deliverable: "Fully managed SaaS platform live"

PHASE 4: Community & Scale (Weeks 10-13) GROWTH SPRINT

Week 10: Agent template marketplace
 Week 11: Community onboarding, Discord launch
 Week 12-13: Bug fixes, performance tuning, hiring
 Deliverable: "1k users, 500+ stars, 10k agents deployed"

SUCCESS CRITERIA:

GitHub: 500+ stars, 100+ forks
Users: 500+ sign-ups (free), 50+ Pro tier
Agents: 10k+ deployed agents in production
Press: Featured in Dev.to, HackerNews, ProductHunt
Revenue: \$5k MRR (Pro + Beta Enterprise)

SLIDE 11: Key Performance Indicators (KPIs)

Metrics to Track & Optimize

PRODUCT METRICS

User Activation Rate: 60%+ of sign-ups create agent
Deploy Success Rate: 98%+ of agents deploy without error
Agent Uptime: 99.8% SLA maintained
Avg Deploy Time: < 2 min (UI → Live in K8s)
Cost per Agent Run: Avg \$0.02-\$0.05

COMMUNITY METRICS

GitHub Stars: Month 1: 100, Month 3: 250, Month 6: 500+
Forks: Month 1: 10, Month 3: 50, Month 6: 150+
Community PRs: Month 1: 5, Month 3: 30, Month 6: 100+
Templates Created: Month 1: 5, Month 3: 50, Month 6: 200
Discord Members: Target 500+ by month 6

BUSINESS METRICS

Sign-ups: Week 1: 100, Month 1: 500, Month 3: 2k+
Free → Pro Conversion: Target 10% (50 Pro users/month 1)
MRR (Monthly Recurring Revenue):

- Month 1: \$500 (early Pro users)
- Month 3: \$5k (50 Pro + 1 Enterprise customer)
- Month 6: \$25k (150 Pro + 3 Enterprise)

Churn Rate: <5% (Pro tier), target: 2%
Lifetime Value (LTV): Pro user avg \$1k/year

ENGAGEMENT METRICS

DAU (Daily Active Users): Target 50% of Pro users
Agents Deployed per User: Target avg 5 agents/user
Avg Runs per Agent per Week: Target 20+ runs
Support Response Time: <4 hours for Pro tier

SLIDE 12: Competitive Landscape & Positioning

Where AgentForge Fits in the Market

COMPETITOR ANALYSIS:

LangChain (Open-Source LLM Framework)

Strengths: Massive community, good APIs

Weaknesses: No UI, no cloud deployment, DIY everything

AgentForge Edge: Visual builder + managed cloud

Devin AI (Commercial AI Coding Agent)

Strengths: Impressive single agent, well-funded

Weaknesses: Closed-source, vendor lock-in, expensive

AgentForge Edge: Open-source, multi-agent, affordable

CrewAI (Multi-Agent Framework)

Strengths: Good multi-agent patterns

Weaknesses: Python-only, no UI, no cloud infra

AgentForge Edge: Visual UI + K8s + observability

Zapier / Make (Low-Code Automation)

Strengths: Mature, huge template library

Weaknesses: Not ML-focused, expensive for agents

AgentForge Edge: AI-first, cheaper for ML workloads

AGENTFORGE POSITIONING:

"Open-source, cloud-native alternative to Devin/DevOps tools"

"Drag-and-drop agent builder with production infra included"

"From prototype to production in 5 minutes, not 5 weeks"

TARGET MARKET: Startups + Enterprise ML teams + Individual
developers building portfolios

SLIDE 13: Call to Action & Contact

NEXT STEPS & GET INVOLVED

LIVE DEMO: <https://demo.agentforge.dev>

GITHUB: <https://github.com/AdityaShenvi/AgentForge>

DOCS: <https://docs.agentforge.dev>

DISCORD: <https://discord.gg/agentforge>

CONTRIBUTE: We're looking for:

- Backend engineers (FastAPI, LangGraph, Kubernetes)
- Frontend engineers (Next.js, React Flow)
- DevOps engineers (Terraform, ArgoCD, observability)
- ML engineers (agent optimization, benchmarking)

CONTACT:

Aditya Shenvi
aditya@agentforge.dev
<https://www.adityacuz.dev>
LinkedIn: /in/adityashenvi

© 2025-26 Aditya Shenvi. Released under MIT License.

"Deploy AI Agents. Change the Future."

END OF PITCH DECK

Watermark: Aditya Shenvi © 2025-26

Document Version: 1.0 / Professional Presentation / December 6, 2025